

CPU:

Architecture: x86_64
CPU(s): 80
Vendor ID: GenuineIntel
Model name: Intel(R) Xeon(R) Gold 6248 CPU @ 2.50GHz
Caches (sum of all):
L1d: 1.3 MiB (40 instances)
L1i: 1.3 MiB (40 instances)
L2: 40 MiB (40 instances)
L3: 55 MiB (2 instances)
NUMA:
NUMA node0 CPU(s): 0-19,40-59
NUMA node1 CPU(s): 20-39,60-79

Название сервера:

ProLiant XL270d Gen10

Доступные numa nodes:

available: 2 nodes (0-1)
node 0 size: 385636 MB
node 1 size: 387008 MB

Система: Ubuntu 22.04.5 LTS

The graph illustrates the performance of three different methods in terms of the number of iterations required as the number of nodes increases from 0 to 40. The 'Linear' method (yellow line) shows a steep, linear increase in iterations, reaching 40 iterations at 40 nodes. The '20000x20000' method (blue line) and the '40000x40000' method (red line) show a much slower increase, with the '40000x40000' method generally requiring fewer iterations than the '20000x20000' method for a given number of nodes.

Nodes	20000x20000	40000x40000	Linear
0	1	2	2
5	4	3	5
10	8	6	10
15	11	10	15
20	12	12	20
25	13	13	25
30	14	13	30
35	15	14	35
40	16	14	40